

A vertical bar on the left side of the slide, composed of several overlapping oval shapes in shades of grey, black, and red.

Collocations

HAP/LAP. Corpus Linguistics.

Frequency lists and beyond



- Key insight: frequency lists are very helpful.
- **But:** top 10 frequent words in any corpora are the same.
- Exploit frequency data to get deeper insights.
 - Collocations.

- What are collocations?
 - systematic co-occurrence of words in use.
 - words which are likely to occur in the context of another.
- Usually within a window of words.

- Collocation is a type of Multi-Word Expression (MLWE)
- What is a MLWE?
 - Lexical unit larger than a word
 - Include the range of phenomena
 - Many terms (with slightly different semantics)
 - *chunk, cliché, collocation, extended lexical unit, fixed expression, formulaic sequence, idiom, idiomatic expression, lexical/lexicalized phrase, multi-word unit, phraseme, phraseologism, phraseological unit, phrasal lexical item, phrasal lexeme, prefabricated chunk, prefab*

- Many types of MLWEs
 - Compounds: “disk drive”
 - Phrasal verbs: “make up”
 - Other phrases: “bacon and eggs”
- May be several words long
 - “international best practice”
- May be discontinuous:
 - “make [something] up”

- MLWEs are often not *compositional*
 - An NLP expression is *compositional* if the meaning of the expression can be predicted from the meaning of the parts.
- Degrees of compositionality:
 - “fish and chips” (collocation): fully compositional, non-idiomatic.
 - “strong tea” (collocation) but *strong* used in a slightly different manner.
 - “black market”: one or several words are idiomatic.
 - “red herring”, “kick the bucket” (idioms): fully opaque, non-compositional.

- **Non-substitutability:** We cannot substitute other words for the components of a MLWE, even if they have the same meaning in the context.
 - “Fast foot” \Rightarrow “Quick food” ?
 - “White wine” \Rightarrow “Yellow wine” ?
- **Non-modifiability:** many MLWEs can not be freely modified with additional lexical material or through grammatical transformations.
 - “The cat has got your tongue” \Rightarrow
“The large, furry cat has your tongue” or
“The cats have your tongues”

- Overlap between MLWEs and notions like
 - *term*
 - *technical term*
 - *terminological term*
- Usually, used when collocations are extracted from technical domains
 - terminology extraction

- Statistical NLP (such as SMT)
 - word translates differently according to MLWE it occurs in
- Information Retrieval (IR)
 - index only “interesting” phrases
 - language models
- Lexicographers
 - frequent ways a word is used.
 - multiword detection (and inclusion in dictionaries)

- Simple case: how to identify **contiguous, two words collocations**.
- First idea: find the most common two word sequences in text (**bigrams**)

$C(w^1, w^2)$	w^1	w^2
80,871	of	the
58,841	in	the
26,430	to	the
21,842	on	the
...

- The table (taken from Manning and Shütze, 1999) shows bigram frequency counts.
- Do not capture the collocations present in the text.
- Frequent bigrams represent common syntactic constructions.
- Many function words: grammatical words which always occur in sentences.

Obtaining bigrams (bigrams.py)

Obtain bigrams from gutenber corpus. Create a program to output bigrams:

- First line is total counts:
N_bigrams TAB *N_tokens*
- Then, one line per bigram:
wordA TAB *wordB* TAB *freq_big* TAB *freq_A* TAB *freq_B*

Example:

2447758	2543994		
, and	41331	192338	78770
of the	18911	70040	125730
in the	9793	31874	125730
; and	7589	27837	78770
and the	6432	78770	125730

- **Problem:** Frequent bigrams give almost no information
- Workarounds:
 1. Take into account individual word frequencies. Measure association by chance.
 2. Filter collocations according to external factor.
 - Use POS tags to get, for instance, *adjective noun*, *noun noun*, etc.

- There are many techniques to overcome the limitations of frequency.
- Hypothesis testing: whether two words occur more frequently than by chance
 - *t* test
 - chi square
 - likelihood ratios
- Mutual information (pointwise mutual information, *pmi*)
 - information theoretically motivated.
 - measure of the variable's mutual dependence.
- Dice coefficient

Contingency table

	$V = v$	$V \neq v$
$U = u$	O_{11}	O_{12}
$U \neq u$	O_{21}	O_{22}

- Count how many times the words u and v appear together,
- and also how many times each word occur by its own.
- For instance, for the words “box” and “black”

	$V = \text{box}$	$V \neq \text{box}$
$U = \text{black}$	123 <i>black box</i>	13,168 <i>e.g. black house</i>
$U \neq \text{black}$	1,810 <i>e.g. white box</i>	4,951,883 <i>e.g. white house</i>

- We can derive many association measures using this matrix.

Contingency table: marginal frequencies



	$V = box$	$V \neq box$	
$U = black$	O_{11}	O_{12}	f_A
$U \neq black$	O_{21}	O_{22}	$f_{\neq A}$
	f_B	$f_{\neq B}$	N

- N = number of words in the corpus.
- Observed frequencies:
 - $f_A = O_{11} + O_{12}$ $f_{\neq A} = O_{21} + O_{22}$
 - $f_B = O_{11} + O_{21}$ $f_{\neq B} = O_{12} + O_{22}$
- Expected frequencies:
 - $E_{11} = \frac{f_A f_B}{N}$ $E_{12} = \frac{f_A f_{\neq B}}{N}$
 - $E_{21} = \frac{f_{\neq A} f_B}{N}$ $E_{22} = \frac{f_{\neq A} f_{\neq B}}{N}$

Association measures



Student t test $\frac{O_{11} - E_{11}}{\sqrt{O_{11}}}$

chi square $\frac{N(O_{11} - E_{11})^2}{E_{11}E_{22}}$ $\frac{N(f_{AB} - \frac{f_A f_B}{N})^2}{f_A f_B f_{\neq A} f_{\neq B}}$

log likelihood $2 \sum_{ij} O_{ij} \log \frac{O_{ij}}{E_{ij}}$ omit $O_{ij} = 0$ from formula

pmi $\log \frac{p(A,B)}{p(A)p(B)} = \log \frac{O_{11}}{E_{11}} = \log(N \frac{f_{AB}}{f_A f_B})$

Dice $\frac{2O_{11}}{f_A + f_B} = \frac{2f_{AB}}{f_A + f_B}$

Jaccard $\frac{O_{11}}{O_{11} + O_{12} + O_{21}}$

Exercise (pmi.py)

Using bigram file, calculate *pmi* association measures. Discard bigrams whose frequency **is below 3**.

Additional issues



- Normalization: case insensitive, etc.
- Filter out non interesting words:
 - Named Entities, numbers, etc.

- PMI for trigrams (as Perl NSP package)

$$pmi(w_1, w_2, w_3) = 2 \log N_t + \log O_{111} - \log O_{1pp} - \log O_{p1p} - \log O_{pp1}$$

where

- N_t : total number of trigrams.
- O_{111} : (w_1, w_2, w_3) trigram count.
- O_{1pp} : number of trigrams starting with w_1 .
- O_{p1p} : number of trigrams starting with w_2 .
- O_{pp1} : number of trigrams starting with w_3 .

Collocations: filter by POS



Tag Pattern	Example
A N	<i>linear function</i>
N N	<i>regression coefficients</i>
A A N	<i>Gaussian random variable</i>
A N N	<i>cumulative distribution function</i>
N A N	<i>mean squared error</i>
N N N	<i>class probability function</i>
N P N	<i>degrees of freedom</i>

- (Justeson and Katz, 1995): pass candidate phrases through a POS filter that are likely to be phrases. See patterns above.

Collocations: filter by POS



$C(w^1, w^2)$	w^1	w^2	Tag Pattern
11487	New	York	A N
7261	United	States	A N
5412	Los	Angeles	N N
3301	last	year	A N
3191	Saudi	Arabia	N N
2699	last	week	A N
2514	vice	president	A N
2378	Persian	Gulf	A N
2161	San	Francisco	N N
...

- The table shows the most highly ranked phrases after applying the filter.
- Surprisingly good results: only two are compositional (“last year”, “last week”)

Exercise (`bigrams_pos.py`)

Obtain most frequent Adj Noun and Noun Noun bigrams from gutenber corpus.