

A vertical bar on the left side of the slide, composed of several overlapping, vertically-oriented oval shapes in shades of grey, black, and red.

Keywords

HAP/LAP. Corpus Linguistics.

- Locate and define words that describe document topics.
- Useful in many areas
 - Search engines
 - Document information
 - Automatic summarization
- Challenging problem due to non regular nature of language.
- Example:

<http://www.cortical.io/demos.html>

Keyword Extraction



- To extract keywords we need more than one corpus
 - Corpus A: the corpus we want to analyze.
 - Corpus B: a reference corpus to compare against.
- Many techniques, based on frequency analysis.

Keyword Extraction: tf-idf



- Basic technique, but gets very good results.
- In principle, frequent terms are candidates to be keywords
 - but some terms are always very frequent (“the”, “of”, ...)
- Idea: analyze frequency of terms within document and across documents.
 - (+) term appears frequently in the document.
 - (-) term appears frequently in all the documents.
- Ideally, we want terms that appear frequently in one document but do not appear in other documents.

- **tf**: term frequency

$\text{tf}(t, d) = f_d(t)$, frequency of term t in document d

- There are other choices:

$$\text{tf}(t, d) = \begin{cases} \log(f_d(t) + 1) & t \in d \\ 0 & \text{otherwise} \end{cases}$$

$$\text{tf}(t, d) = \frac{1}{2} + \frac{\frac{1}{2}f_d(t)}{\max\{f_d(w) : w \in d\}}$$

- The last formula tries to minimize the fact that longer documents have higher frequencies.

- **idf**: inverse term frequency: how much information provided by the term.

$$\text{idf}(t, D) = \log \frac{N}{|\{d \in D : t \in d\}|}$$

where

t the term

D set of all documents

N number of documents in D

$|\{d \in D : t \in d\}|$ number of documents that contain term t

Keyword Extraction: tf-idf



- Putting tf and idf together:

$$\text{tfidf}(t, d, D) = \text{tf}(t, d) \times \text{idf}(t, D)$$

- Alternative to tf/idf

- $$\text{tf}(t, d) = \frac{f_d(t) \cdot (k_1 + 1)}{f_d(t) + k_1 \cdot \left(1 - b + b \cdot \frac{|d|}{\text{avgdl}}\right)}$$

where

$|d|$ length of document.

avgdl: average length of documents.

k_1 : free parameters (usually $k_1 \in [1.2, 2.0]$)

b (usually $b = 0.75$)

- $$\text{idf}(t, D) = \log \frac{N - n(t) + 0.5}{n(t) + 0.5}$$

N size of D

$n(t)$ number of documents containing t

- More information: <https://labur.eus/sc069>


Exercise

- Using `df.py`, create `word.df` file with following format:
Number_of_documents
word TAB number_of_documents_where_the_word_appears
- Complete `tfidf.py` script that given `word.df` and a `.tab` file outputs the `tfidf` values of words:
word TAB `tfidf`
- Note: use the formula $tfidf = \log(tf + 1) \times idf$

Vector Space model



- Represent documents as vectors
- Dimensions are vocabulary words
- Document similarity is proportional to the angle between vectors:
 - parallel: documents are the same
 - perpendicular: documents are completely different



G. Salton , A. Wong , C. S. Yang, "A vector space model for automatic indexing, Communications of the ACM", v.18 n.11, p.613-620, Nov. 1975

Using keywords to compare documents



Idea: use keywords to compare documents against.

- Create a vector for each document.

- topK keywords according to tf-idf

$$d_1 = \mathbf{u} = (w_1 : \text{tfidf}(w_1, d_1, D), w_2 : \text{tfidf}(w_2, d_1, D), \dots)$$

$$d_2 = \mathbf{v} = (w_4 : \text{tfidf}(w_4, d_2, D), w_{10} : \text{tfidf}(w_{10}, d_2, D), \dots)$$

Document	whale	bride	widow	bullet	...
Moby Dick	10.1	2.0	1.4	0.9	...
Emma	0.0	3.4	0.9	0.0	...
Father Brown	0.0	0.0	0.5	5.1	...
...

Exercise

- Check `docSimilarity.py` to obtain similarity scores of two documents.

$$\cos(d_j, q) = \frac{\mathbf{d}_j \cdot \mathbf{q}}{\|\mathbf{d}_j\| \|\mathbf{q}\|} = \frac{\sum_{i=1}^N w_{i,j} w_{i,q}}{\sqrt{\sum_{i=1}^N w_{i,j}^2} \sqrt{\sum_{i=1}^N w_{i,q}^2}}$$

Exercise

- Obtain similarity among documents using list of terms weighted by tf-idf. Use `create_tfidf.sh` script for creating all `tfidf` vectors, and the `docsim.csv` file.
- Open `docsim.csv` in libreoffice calc and compare documents.